

Problem 1

1a) False, 0.8 applies to the logistic cdf, and since it is non-linear the effect cannot just be estimated through the coefficient, one would have to take the derivative of the logistic cdf.

1b) Note below lowercase phi represents the pdf of the zero centered gaussian distribution with variance 1.

$$\frac{dP}{d(\text{age})} = \frac{d(\Phi(\alpha + \beta \cdot \text{age} + \gamma \text{age}^2 + \delta \text{educ}))}{d(\text{age})}$$

$$\frac{dP}{d(\text{age})} = \frac{d(\Phi(\alpha + \beta \cdot \text{age} + \gamma \text{age}^2 + \delta \text{educ}))}{d(\alpha + \beta \cdot \text{age} + \gamma \text{age}^2 + \delta \text{educ})} \cdot \frac{\Phi(\alpha + \beta \cdot \text{age} + \gamma \text{age}^2 + \delta \text{educ})}{d(\text{age})}$$

$$\frac{dP}{d(\text{age})} = \phi(\alpha + \beta \cdot \text{age} + \gamma \text{age}^2 + \delta \text{educ}) \cdot (\beta \cdot 1 + 2 \cdot \delta \cdot \text{age})$$

1c) Yes potentially, you would be able to use this to determine whether we can utilize a distribution that favors the heavy tails more or not. If you have heavy tails then you may want to prefer using a normal cdf.

Problem 2

```
In [15]: import pandas as pd
import statsmodels.api as sm
import statsmodels.discrete as smd

df = pd.read_stata('JTRAIN2.dta')
```

2a) Run a linear regression of train on several demographic and pretraining variables:
 unem74,unem75,age,educ,black,hisp,married. Are these variables jointly significant at the 5% level

```
In [16]: Y = df['train']
X = sm.add_constant(df[['unem74', 'unem75', 'age', 'educ', 'black', 'hisp', 'married'])

model = sm.OLS(exog=X, endog=Y).fit()

model.summary()
```

Out[16]:

OLS Regression Results

Dep. Variable:	train	R-squared:	0.022			
Model:	OLS	Adj. R-squared:	0.007			
Method:	Least Squares		F-statistic: 1.429			
Date:	Mon, 03 Nov 2025	Prob (F-statistic):	0.192			
Time:	09:17:41	Log-Likelihood:	-311.53			
No. Observations:	445	AIC:	639.1			
Df Residuals:	437	BIC:	671.8			
Df Model:	7					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
const	0.3380	0.189	1.784	0.075	-0.034	0.710
unem74	0.0209	0.077	0.270	0.787	-0.131	0.173
unem75	-0.0956	0.072	-1.329	0.184	-0.237	0.046
age	0.0032	0.003	0.942	0.347	-0.003	0.010
educ	0.0120	0.013	0.900	0.368	-0.014	0.038
black	-0.0817	0.088	-0.931	0.352	-0.254	0.091
hisp	-0.2000	0.117	-1.710	0.088	-0.430	0.030
married	0.0373	0.064	0.579	0.563	-0.089	0.164
	Omnibus:	2209.423	Durbin-Watson:	0.035		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	68.291			
Skew:	0.327	Prob(JB):	1.48e-15			
Kurtosis:	1.195	Cond. No.	250.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [17]:

```
print("X"*100)
print("Model is not jointly significant at 5 percent level since p-val is 0.1")
print("X"*100)
```

XX
 XXXXXXXXXXXXXXXXXXXXXXX

Model is not jointly significant at 5 percent level since p-val is 0.192

XX
 XXXXXXXXXXXXXXXXXXXXXXX

```
In [18]: Y = df['train']
X = sm.add_constant(df[['unem74', 'unem75', 'age', 'educ', 'black', 'hisp', 'married']])

probit_model = smd.discrete_model.Probit(exog=X, endog=Y).fit()
probit_model.summary()
```

Optimization terminated successfully.
 Current function value: 0.667436
 Iterations 5

Out[18]: Probit Regression Results

Dep. Variable:	train	No. Observations:	445			
Model:	Probit	Df Residuals:	437			
Method:	MLE	Df Model:	7			
Date:	Mon, 03 Nov 2025	Pseudo R-squ.:	0.01685			
Time:	09:17:41	Log-Likelihood:	-297.01			
converged:	True	LL-Null:	-302.10			
Covariance Type:	nonrobust	LLR p-value:	0.1785			
	coef	std err	z	P> z	[0.025	0.975]
const	-0.4241	0.487	-0.871	0.384	-1.379	0.530
unem74	0.0530	0.199	0.266	0.790	-0.338	0.444
unem75	-0.2477	0.185	-1.339	0.181	-0.610	0.115
age	0.0083	0.009	0.948	0.343	-0.009	0.026
educ	0.0314	0.034	0.916	0.360	-0.036	0.099
black	-0.2069	0.225	-0.920	0.358	-0.648	0.234
hisp	-0.5398	0.309	-1.750	0.080	-1.144	0.065
married	0.0966	0.166	0.584	0.560	-0.228	0.421

```
In [19]: print("X"*100)
print(f"The ratio test p-value is 0.1785 meaning that the model is not significant")
print("X"*100)
```

XX
 XXXXXXXXXXXXXXXXXXXXXXX

The ratio test p-value is 0.1785 meaning that the model is not significant
 XXX
 XXXXXXXXXXXXXXXXXXXXXXX

Based on your answers to parts (a) and (b), does it appear that participation in job training can be treated as exogenous for explaining 1978 unemployment status? Explain

Yes it can be because the other variables aren't jointly significant (and therefore they are not considered a multivariate combination of previous data). It can be seen as exogenous for explaining 1978 unemployment status.

d) Simple OLS regression of unem78 train and report the results in equation form.

```
In [20]: Y = df['unem78']
X = sm.add_constant(df['train'])

model = sm.OLS(exog=X, endog=Y).fit()

model.summary()
```

Out[20]: OLS Regression Results

Dep. Variable:	unem78	R-squared:	0.014			
Model:	OLS	Adj. R-squared:	0.012			
Method:	Least Squares	F-statistic:	6.265			
Date:	Mon, 03 Nov 2025	Prob (F-statistic):	0.0127			
Time:	09:17:41	Log-Likelihood:	-284.30			
No. Observations:	445	AIC:	572.6			
Df Residuals:	443	BIC:	580.8			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	0.3538	0.028	12.419	0.000	0.298	0.410
train	-0.1106	0.044	-2.503	0.013	-0.197	-0.024
Omnibus:	529.053	Durbin-Watson:	2.008			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	79.145			
Skew:	0.813	Prob(JB):	6.51e-18			
Kurtosis:	1.727	Cond. No.	2.47			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [21]: print("X"*100)
print("Since p-value on train is 0.013, this is very statistically signficia
```

```
print("X"*100)
```

```
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
XXXXXXXXXXXXXXXXXXXXXXXXXXXX
Since p-value on train is 0.013, this is very statistically significant.
XXXXXXXXXXXXXXXXXXXXXXXXXXXX
XXXXXXXXXXXXXXXXXXXXXXXXXXXX
```

e) Run a probit of unem78 on train and report the results in equation form. Does it make sense to compare the probit coefficient on train with the coefficient obtained from the linear model in part (d)?

```
In [22]: Y = df['unem78']
X = sm.add_constant(df['train'])

probit_model = smf.discrete_model.Probit(exog=X, endog=Y).fit()
probit_model.summary()
```

```
Optimization terminated successfully.
    Current function value: 0.610298
    Iterations 5
```

```
Out[22]: Probit Regression Results

Dep. Variable: unem78 No. Observations: 445
Model: Probit Df Residuals: 443
Method: MLE Df Model: 1
Date: Mon, 03 Nov 2025 Pseudo R-squ.: 0.01147
Time: 09:17:41 Log-Likelihood: -271.58
converged: True LL-Null: -274.73
Covariance Type: nonrobust LLR p-value: 0.01204

      coef  std err      z   P>|z|  [0.025  0.975]
const -0.3750  0.080  -4.702  0.000   -0.531  -0.219
train -0.3210  0.128  -2.498  0.012   -0.573  -0.069
```

You can only really compare the sign of the effects in the OLS regression and the probit model. Both show statistically significant relationships. It doesn't make sense to compare them in magnitude however

(f) Find the fitted probabilities from parts (d) and (e). Explain why they are identical.

```
In [23]: import numpy as np

OLS_preds = list(model.predict())
probit_preds = list(probit_model.predict())
```

```
# Check if predictions are approximately equal (within default tolerance)
if np.allclose(OLS_preds, probit_preds):
    print("They are approximately the same")
else:
    print("Not the same")

# Show first 10 predictions from each model
print("\nFirst 10 OLS predictions:", OLS_preds[:10])
print("First 10 Probit predictions:", probit_preds[:10])

# Show maximum difference
max_diff = np.max(np.abs(np.array(OLS_preds) - np.array(probit_preds)))
print(f"\nMaximum absolute difference: {max_diff:.10f}")
```

They are approximately the same

```
First 10 OLS predictions: [0.24324324324324356, 0.24324324324324356, 0.24324324324356, 0.24324324324356, 0.24324324324356, 0.24324324324356, 0.24324324324356, 0.24324324324356, 0.24324324324356, 0.24324324324356]
First 10 Probit predictions: [0.24324324324324326, 0.24324324324324326, 0.24324324324324326, 0.24324324324324326, 0.24324324324324326, 0.24324324324324326, 0.24324324324324326, 0.24324324324324326, 0.24324324324324326, 0.24324324324324326]
```

Maximum absolute difference: 0.0000000000

The predictions are the same because when you do univariate OLS and univariate probit model, the optimization must be the same numerically in terms of the maximum of likelihood at standard OLS procedure.

(g) Add all the controls from part (a) both to the linear regression and to the probit regression of unem78 on train. Are the fitted probabilities now identical?

```
In [24]: # OLS Model: unem78 on train + all controls
Y = df['unem78']
X = sm.add_constant(df[['train', 'unem74', 'unem75', 'age', 'educ', 'black', 'his'])

model_OLS = sm.OLS(endog=Y, exog=X).fit()

print("=*80")
print("OLS Model")
print("=*80")
print(model_OLS.summary())
print("\n")

# Probit Model: unem78 on train + all controls
Y = df['unem78']
X = sm.add_constant(df[['train', 'unem74', 'unem75', 'age', 'educ', 'black', 'his'])

probit_model = smm.discrete_model.Probit(endog=Y, exog=X).fit()
```

```
print("=*80")
print("Probit Model")
print("=*80")
print(probit_model.summary())
print("\n")

# Compare predictions
OLS_preds = list(model_OLS.predict())
probit_preds = list(probit_model.predict())

# Check if predictions are approximately equal (within default tolerance)
if np.allclose(OLS_preds, probit_preds):
    print("They are approximately the same")
else:
    print("Not the same")

# Show first 10 predictions from each model
print("\nFirst 10 OLS predictions:", OLS_preds[:10])
print("First 10 Probit predictions:", probit_preds[:10])

# Show maximum difference
max_diff = np.max(np.abs(np.array(OLS_preds) - np.array(probit_preds)))
print(f"\nMaximum absolute difference: {max_diff:.10f}")

print("\n" + "=*80")
print("Answer: No, the fitted probabilities are NOT identical when we includ")
print("multiple covariates. Unlike the univariate case in part (f), the OLS")
print("probit models produce different predictions when controls are added.")
print("=*80")
```

OLS Regression Results						
Dep. Variable:	unem78	R-squared:	0.0			
Model:	OLS	Adj. R-squared:	0.0			
Date:	Mon, 03 Nov 2025	Prob (F-statistic):	0.007			
No. Observations:	445	AIC:	57			
Df Residuals:	436	BIC:	60			
Df Model:	8					
Covariance Type:	nonrobust					
coef	std err	t	P> t	[0.025	0.97	
const	0.1632	0.176	0.927	0.355	-0.183	0.5
train	-0.1117	0.044	-2.521	0.012	-0.199	-0.0
unem74	0.0387	0.072	0.540	0.589	-0.102	0.1
unem75	0.0160	0.067	0.239	0.811	-0.115	0.1
age	4.332e-05	0.003	0.014	0.989	-0.006	0.0
educ	0.0001	0.012	0.012	0.991	-0.024	0.0
black	0.1888	0.081	2.322	0.021	0.029	0.3
hispanic	-0.0377	0.109	-0.347	0.729	-0.251	0.1
married	-0.0254	0.060	-0.426	0.670	-0.143	0.0
Omnibus:	440.373	Durbin-Watson:	2.0			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	70.7			
Skew:	0.749	Prob(JB):	4.41e-			
16						

Kurtosis: 1.748 Cond. No. 25

1.

=====

==

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Optimization terminated successfully.

Current function value: 0.591714

Iterations 5

=====

==

Probit Model

=====

====

Probit Regression Results

=====

==

Dep. Variable: unem78 No. Observations: 4

45

Model: Probit Df Residuals: 4

36

Method: MLE Df Model: 4

8

Date: Mon, 03 Nov 2025 Pseudo R-squ.: 0.041

58

Time: 09:17:41 Log-Likelihood: -263.

31

converged: True LL-Null: -274.

73

Covariance Type: nonrobust LLR p-value: 0.0035

70

=====

==

	coef	std err	z	P> z	[0.025	0.97
5]						

--

const	-1.0103	0.538	-1.878	0.060	-2.065	0.0
-------	---------	-------	--------	-------	--------	-----

44

train	-0.3366	0.132	-2.557	0.011	-0.595	-0.0
-------	---------	-------	--------	-------	--------	------

79

unem74	0.1061	0.213	0.499	0.618	-0.311	0.5
--------	--------	-------	-------	-------	--------	-----

23

unem75	0.0636	0.197	0.323	0.747	-0.323	0.4
--------	--------	-------	-------	-------	--------	-----

50

age	0.0007	0.009	0.074	0.941	-0.017	0.0
-----	--------	-------	-------	-------	--------	-----

19

educ	-0.0019	0.037	-0.051	0.959	-0.074	0.0
------	---------	-------	--------	-------	--------	-----

70

black	0.6337	0.274	2.310	0.021	0.096	1.1
-------	--------	-------	-------	-------	-------	-----

71

hisp	-0.1649	0.379	-0.435	0.663	-0.908	0.5
------	---------	-------	--------	-------	--------	-----

```

78
married      -0.0778      0.177      -0.439      0.661      -0.425      0.2
69
=====
==
```

Not the same

First 10 OLS predictions: [0.27271825742439904, 0.07068344139995397, 0.29799
65719925336, 0.29772237640252935, 0.29754955547604456, 0.2972173029454226,
0.2976933479322205, 0.2979389650170387, 0.29822699989451334, 0.0838564347934
373]

First 10 Probit predictions: [0.26857823653209445, 0.08942352721999941, 0.29
2542461326591, 0.2924958423282561, 0.29584836731645947, 0.2926348585488725,
0.2909187965580752, 0.2936580891182099, 0.28809976788650393, 0.1046700703145
2641]

Maximum absolute difference: 0.0643153572

```

=====
=====
Answer: No, the fitted probabilities are NOT identical when we include  
multiple covariates. Unlike the univariate case in part (f), the OLS and  
probit models produce different predictions when controls are added.
```

```

=====
=====
          OLS Regression Results
=====
=====
Dep. Variable:           unem78    R-squared:         0.0
46
Model:                 OLS     Adj. R-squared:      0.0
29
Method:                Least Squares   F-statistic:       2.6
40
Date:                  Mon, 03 Nov 2025   Prob (F-statistic):  0.007
80
Time:                  09:17:41        Log-Likelihood:   -276.
90
No. Observations:      445        AIC:             57
1.8
Df Residuals:          436        BIC:             60
8.7
Df Model:               8
Covariance Type:       nonrobust
=====
```

```

=====
      coef    std err        t     P>|t|    [0.025    0.97
5]
-----
const      0.1632      0.176      0.927      0.355      -0.183      0.5
09
train     -0.1117      0.044     -2.521      0.012      -0.199      -0.0
25
```

unem74	0.0387	0.072	0.540	0.589	-0.102	0.1
79						
unem75	0.0160	0.067	0.239	0.811	-0.115	0.1
47						
age	4.332e-05	0.003	0.014	0.989	-0.006	0.0
06						
educ	0.0001	0.012	0.012	0.991	-0.024	0.0
24						
black	0.1888	0.081	2.322	0.021	0.029	0.3
49						
hisp	-0.0377	0.109	-0.347	0.729	-0.251	0.1
76						
married	-0.0254	0.060	-0.426	0.670	-0.143	0.0
92						
<hr/>						
==						
Omnibus:	440.373	Durbin-Watson:	2.0			
04						
Prob(Omnibus):	0.000	Jarque-Bera (JB):	70.7			
14						
Skew:	0.749	Prob(JB):	4.41e-			
16						
Kurtosis:	1.748	Cond. No.	25			
1.						
<hr/>						
==						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Optimization terminated successfully.

Current function value: 0.591714

Iterations 5

<hr/>						
====						
Probit Model						
<hr/>						
====						
Probit Regression Results						
<hr/>						
==						
Dep. Variable:	unem78	No. Observations:	4			
45						
Model:	Probit	Df Residuals:	4			
36						
Method:	MLE	Df Model:				
8						
Date:	Mon, 03 Nov 2025	Pseudo R-squ.:	0.041			
58						
Time:	09:17:41	Log-Likelihood:	-263.			
31						
converged:	True	LL-Null:	-274.			
73						
Covariance Type:	nonrobust	LLR p-value:	0.0035			

	coef	std err	z	P> z	[0.025	0.97
5]						
--						
const	-1.0103	0.538	-1.878	0.060	-2.065	0.0
44						
train	-0.3366	0.132	-2.557	0.011	-0.595	-0.0
79						
unem74	0.1061	0.213	0.499	0.618	-0.311	0.5
23						
unem75	0.0636	0.197	0.323	0.747	-0.323	0.4
50						
age	0.0007	0.009	0.074	0.941	-0.017	0.0
19						
educ	-0.0019	0.037	-0.051	0.959	-0.074	0.0
70						
black	0.6337	0.274	2.310	0.021	0.096	1.1
71						
hisp	-0.1649	0.379	-0.435	0.663	-0.908	0.5
78						
married	-0.0778	0.177	-0.439	0.661	-0.425	0.2
69						
--						
==						

Not the same

First 10 OLS predictions: [0.27271825742439904, 0.07068344139995397, 0.2979965719925336, 0.29772237640252935, 0.29754955547604456, 0.2972173029454226, 0.2976933479322205, 0.2979389650170387, 0.29822699989451334, 0.0838564347934373]

First 10 Probit predictions: [0.26857823653209445, 0.08942352721999941, 0.292542461326591, 0.2924958423282561, 0.29584836731645947, 0.2926348585488725, 0.2909187965580752, 0.2936580891182099, 0.28809976788650393, 0.10467007031452641]

Maximum absolute difference: 0.0643153572

=====

====

Answer: No, the fitted probabilities are NOT identical when we include multiple covariates. Unlike the univariate case in part (f), the OLS and probit models produce different predictions when controls are added.

=====

====

Clearly the predictions are not the same thanks!

Calculate Average Partial Effects (APE) for both Linear and Probit Models

We need to compute the APE for all variables in both models and compare them.

```
In [25]: # Calculate Average Partial Effects (APE) from the models in part (g)

from scipy.stats import norm

# For Linear Model (OLS): APE = coefficient (constant across all observations)
ape_ols = model_OLS.params

print("=*80")
print("Average Partial Effects - LINEAR MODEL (OLS)")
print("=*80")
print("\nFor the linear probability model, the partial effect is constant and")
print("equals the coefficient for each variable:\n")
print(ape_ols)
print("\n")

# For Probit Model: APE = mean( $\varphi(X'\beta)$  *  $\beta_j$ ) for each variable j
# where  $\varphi$  is the standard normal PDF

# Recreate X matrix for predictions
X = sm.add_constant(df[['train', 'unem74', 'unem75', 'age', 'educ', 'black', 'his'])

# Get the linear predictor  $X'\beta$  for all observations
X_beta = probit_model.predict(X, which='linear')

# Calculate  $\varphi(X'\beta)$  - the standard normal PDF evaluated at  $X'\beta$ 
phi_xbeta = norm.pdf(X_beta)

# Calculate APE for each variable
# APE_j = mean( $\varphi(X'\beta)$ ) *  $\beta_j$ 
mean_phi = phi_xbeta.mean()
ape_probit = mean_phi * probit_model.params

print("=*80")
print("Average Partial Effects - PROBIT MODEL")
print("=*80")
print("\nFor the probit model, APE = mean( $\varphi(X'\beta)$ ) *  $\beta_j$ ")
print(f"Mean( $\varphi(X'\beta)$ ) = {mean_phi:.6f}\n")
print(ape_probit)
print("\n")
```

```
=====
=====
Average Partial Effects - LINEAR MODEL (OLS)
=====
```

For the linear probability model, the partial effect is constant and equals the coefficient for each variable:

```
const      0.163182
train     -0.111703
unem74    0.038693
unem75    0.015961
age       0.000043
educ      0.000144
black     0.188833
hisp      -0.037701
married   -0.025437
dtype: float64
```

```
=====
=====
Average Partial Effects - PROBIT MODEL
=====
```

For the probit model, $APE = \text{mean}(\varphi(X'\beta)) * \beta_j$
 $\text{Mean}(\varphi(X'\beta)) = 0.336235$

```
const      -0.339709
train     -0.113173
unem74    0.035673
unem75    0.021389
age       0.000227
educ      -0.000636
black     0.213062
hisp      -0.055459
married   -0.026148
dtype: float64
```

In [26]: # Compare APE from both models side by side

```
comparison_df = pd.DataFrame({
    'Variable': ape_ols.index,
    'APE - Linear Model': ape_ols.values,
    'APE - Probit Model': ape_probit.values,
    'Ratio (Linear/Probit)': ape_ols.values / ape_probit.values
})

print("*"*80)
print("COMPARISON: Average Partial Effects")
print("*"*80)
print("\n")
```

```

print(comparison_df.to_string(index=False))
print("\n")

# Also show the scaling factor
print(f"Scaling factor (mean  $\phi(X'\beta)$ ): {mean_phi:.6f}")
print(f"\nNote: The probit APE  $\approx$  Linear APE * {mean_phi:.6f}")
print("\n")

# Visual comparison (excluding intercept)
print("=*80")
print("INTERPRETATION")
print("=*80")
print("****")
Key observations:

1. SIGN: Both models give the same sign for all variables (positive or negative)

2. MAGNITUDE: The probit APE values are smaller in absolute value than the linear APE values because they are scaled by  $\phi(X'\beta)$ , which is approximately {:.4f} on average

3. RELATIVE EFFECTS: The ratio of effects between variables is similar in both models

4. For the LINEAR MODEL:
   - The partial effect is CONSTANT across all observations
   -  $APE_j = \beta_j$  (the coefficient itself)

5. For the PROBIT MODEL:
   - The partial effect varies across observations
   -  $APE_j = \text{mean}(\phi(X'\beta)) * \beta_j$ 
   -  $\phi(X'\beta)$  is the standard normal PDF evaluated at the linear predictor

6. The probit model is more appropriate when we want to ensure predicted probabilities stay between 0 and 1, while the linear model can produce predictions outside this range
*****.format(mean_phi))

```

=====
COMPARISON: Average Partial Effects
=====

Variable	APE - Linear Model	APE - Probit Model	Ratio (Linear/Probit)
const	0.163182	-0.339709	-0.480359
train	-0.111703	-0.113173	0.987007
unem74	0.038693	0.035673	1.084660
unem75	0.015961	0.021389	0.746247
age	0.000043	0.000227	0.190653
educ	0.000144	-0.000636	-0.226791
black	0.188833	0.213062	0.886282
hisp	-0.037701	-0.055459	0.679801
married	-0.025437	-0.026148	0.972808

Scaling factor (mean $\varphi(X'\beta)$): 0.336235

Note: The probit APE \approx Linear APE * 0.336235

=====
INTERPRETATION
=====

Key observations:

1. **SIGN:** Both models give the same sign for all variables (positive or negative effect)
2. **MAGNITUDE:** The probit APE values are smaller in absolute value than the linear model
because they are scaled by $\varphi(X'\beta)$, which is approximately 0.3362 on average
3. **RELATIVE EFFECTS:** The ratio of effects between variables is similar in both models
4. **For the LINEAR MODEL:**
 - The partial effect is CONSTANT across all observations
 - $APE_j = \beta_j$ (the coefficient itself)
5. **For the PROBIT MODEL:**
 - The partial effect varies across observations
 - $APE_j = \text{mean}(\varphi(X'\beta)) * \beta_j$
 - $\varphi(X'\beta)$ is the standard normal PDF evaluated at the linear predictor
6. The probit model is more appropriate when we want to ensure predicted probabilities stay between 0 and 1, while the linear model can produce predictions outside this range

```
ide [0,1].
```

```
In [27]: # Create a cleaner comparison table (excluding the intercept)

comparison_clean = pd.DataFrame({
    'Variable': ape_ols.index[1:], # Exclude intercept
    'APE - Linear Model': ape_ols.values[1:],
    'APE - Probit Model': ape_probit.values[1:],
    'Difference': ape_ols.values[1:] - ape_probit.values[1:]
})

print("=*80")
print("AVERAGE PARTIAL EFFECTS COMPARISON (Excluding Intercept)")
print("=*80")
print("\n")
print(comparison_clean.to_string(index=False))
print("\n")

# Summary statistics
print("=*80")
print("SUMMARY")
print("=*80")
print(f"\nMean absolute difference: {np.abs(comparison_clean['Difference'])}")
print(f"Max absolute difference: {np.abs(comparison_clean['Difference'])}.max")
print(f"\nAverage ratio (Linear/Probit): {(ape_ols.values[1:] / ape_probit.v}")
print(f"This is approximately 1/\phi(X'\beta) = 1/{mean_phi:.6f} = {1/mean_phi:.4f}
```

=====
AVERAGE PARTIAL EFFECTS COMPARISON (Excluding Intercept)
=====
=====

Variable	APE - Linear Model	APE - Probit Model	Difference
train	-0.111703	-0.113173	0.001470
unem74	0.038693	0.035673	0.003020
unem75	0.015961	0.021389	-0.005427
age	0.000043	0.000227	-0.000184
educ	0.000144	-0.000636	0.000780
black	0.188833	0.213062	-0.024229
hisp	-0.037701	-0.055459	0.017758
married	-0.025437	-0.026148	0.000711

=====
=====
SUMMARY
=====
=====

Mean absolute difference: 0.006697

Max absolute difference: 0.024229

Average ratio (Linear/Probit): 0.6651

This is approximately $1/\phi(X'\beta) = 1/0.336235 = 2.9741$

Conclusion: Comparing Linear and Probit APE How do they compare? 1. Direction of effects: Both models agree on the direction (sign) of the effect for each variable 2. Magnitude: The probit APE values are consistently smaller in absolute value than the linear model APE because the probit APE is scaled by the standard normal PDF. 3. Scaling relationship: The linear APE is approximately equal to Probit APE divided by the mean of the standard normal PDF evaluated at the linear predictor. 4. Relative importance: The ranking of variables by importance is similar in both models. 5. Interpretation: - Linear model: A one-unit increase in a variable changes the probability by a constant amount (the coefficient) for all observations. - Probit model: A one-unit increase in a variable changes the probability by an amount that varies across observations; the APE is the average of these individual effects. 6. Which to use: The probit model is theoretically superior because it constrains predicted probabilities to [0,1], but the linear model is simpler for interpretation.

(i) Re-estimate regression in (g) as logit, calculate its average partial effect and compare it with your results in (h)

```
In [28]: # Logit Model: unem78 on train + all controls
Y = df['unem78']
X = sm.add_constant(df[['train', 'unem74', 'unem75', 'age', 'educ', 'black', 'hisp']])

logit_model = smf.discrete_model.Logit(endog=Y, exog=X).fit()

print("=*80")
print("LOGIT MODEL")
print("=*80")
print(logit_model.summary())
print("\n")
```

```

Optimization terminated successfully.
  Current function value: 0.591959
  Iterations 6
=====
=====
LOGIT MODEL
=====
=====
      Logit Regression Results
=====
=====
Dep. Variable: unem78   No. Observations: 4
45
Model: Logit   Df Residuals: 4
36
Method: MLE   Df Model: 8
Date: Mon, 03 Nov 2025   Pseudo R-squ.: 0.041
18
Time: 09:17:41   Log-Likelihood: -263.
42
converged: True   LL-Null: -274.
73
Covariance Type: nonrobust   LLR p-value: 0.0038
78
=====
=====
      coef    std err        z     P>|z|    [0.025    0.97
5]
-----
const      -1.7073    0.911   -1.875    0.061   -3.492    0.0
77
train      -0.5532    0.220   -2.516    0.012   -0.984   -0.1
22
unem74      0.1958    0.352    0.556    0.578   -0.495    0.8
86
unem75      0.0827    0.325    0.255    0.799   -0.553    0.7
19
age         0.0004    0.015    0.024    0.981   -0.029    0.0
30
educ       -0.0016    0.061   -0.026    0.979   -0.120    0.1
17
black       1.1024    0.501    2.201    0.028    0.121    2.0
84
hisp        -0.2436    0.694   -0.351    0.725   -1.603    1.1
16
married     -0.1358    0.296   -0.458    0.647   -0.717    0.4
45
=====
=====

```

In [29]: # Calculate Average Partial Effects (APE) for LOGIT Model

```
# For Logit Model: APE = mean(Λ(X'β) * (1 - Λ(X'β)) * β_j) for each variable
# where Λ is the logistic CDF (which equals the predicted probability)

# Get predicted probabilities P(Y=1|X) = Λ(X'β)
predicted_probs = logit_model.predict(X)

# For logit, the marginal effect at each observation is: λ(X'β) * β_j
# where λ(X'β) = Λ(X'β) * (1 - Λ(X'β)) is the logistic PDF
lambda_xbeta = predicted_probs * (1 - predicted_probs)

# Calculate APE for each variable
# APE_j = mean(λ(X'β)) * β_j
mean_lambda = lambda_xbeta.mean()
ape_logit = mean_lambda * logit_model.params

print("=*80")
print("Average Partial Effects – LOGIT MODEL")
print("=*80")
print("\nFor the logit model, APE = mean(Λ(X'β) * (1 - Λ(X'β))) * β_j")
print(f"Mean(λ(X'β)) = Mean(Λ(X'β) * (1 - Λ(X'β))) = {mean_lambda:.6f}\n")
print(ape_logit)
print("\n")
```

```
=====
=====
Average Partial Effects – LOGIT MODEL
=====
=====
```

For the logit model, APE = mean(Λ(X'β) * (1 - Λ(X'β))) * β_j
 Mean(λ(X'β)) = Mean(Λ(X'β) * (1 - Λ(X'β))) = 0.203070

```
const      -0.346708
train      -0.112330
unem74     0.039770
unem75     0.016791
age        0.000073
educ       -0.000321
black      0.223870
hisp       -0.049476
married    -0.027580
dtype: float64
```

In [30]: # Compare APE from all three models: Linear, Probit, and Logit

```
comparison_all = pd.DataFrame({
    'Variable': ape_ols.index,
    'APE - Linear': ape_ols.values,
    'APE - Probit': ape_probit.values,
    'APE - Logit': ape_logit.values,
    'Probit/Linear': ape_probit.values / ape_ols.values,
    'Logit/Linear': ape_logit.values / ape_ols.values,
    'Logit/Probit': ape_logit.values / ape_probit.values
})
```

```
print("=*80")
print("COMPARISON: Average Partial Effects – All Three Models")
print("=*80)
print("\n")
print(comparison_all.to_string(index=False))
print("\n")

print("=*80)
print("SCALING FACTORS")
print("=*80)
print(f"Linear Model: APE = β_j (constant partial effect)")
print(f"Probit Model: Mean(φ(X'β)) = {mean_phi:.6f}")
print(f"Logit Model: Mean(λ(X'β)) = {mean_lambda:.6f}")
print(f"\nRatio: Logit/Probit scaling = {mean_lambda/mean_phi:.4f}")
print("\n")
```

COMPARISON: Average Partial Effects – All Three Models

Variable	APE - Linear ar Logit/Probit	APE - Probit	APE - Logit	Probit/Linear	Logit/Line
const	0.163182	-0.339709	-0.346708	-2.081774	-2.1246
train	1.020604	-0.111703	-0.113173	-0.112330	1.013164
unem74	0.992550	0.038693	0.035673	0.039770	0.921948
unem75	1.114874	0.015961	0.021389	0.016791	1.340039
age	0.785046	0.000043	0.000227	0.000073	5.245132
educ	0.323491	0.000144	-0.000636	-0.000321	-4.409338
black	0.505409	0.188833	0.213062	0.223870	-2.2285
hisp	1.050728	-0.037701	-0.055459	-0.049476	1.128308
married	0.892125	-0.025437	-0.026148	-0.027580	1.471019
	1.054755				1.3123

SCALING FACTORS

Linear Model: APE = β_j (constant partial effect)

Probit Model: Mean($\phi(X'\beta)$) = 0.336235

Logit Model: Mean($\lambda(X'\beta)$) = 0.203070

Ratio: Logit/Probit scaling = 0.6040

COMPARISON: Average Partial Effects – All Three Models

Variable	APE - Linear ar Logit/Probit	APE - Probit	APE - Logit	Probit/Linear	Logit/Line
const	0.163182	-0.339709	-0.346708	-2.081774	-2.1246
train	1.020604				
unem74	-0.111703	-0.113173	-0.112330	1.013164	1.0056
unem75	0.992550				
unem74	0.038693	0.035673	0.039770	0.921948	1.0278
unem75	1.114874				
	0.015961	0.021389	0.016791	1.340039	1.0519

92	0.785046					
	age 0.000043	0.000227	0.000073	5.245132	1.6967	
55	0.323491					
	educ 0.000144	-0.000636	-0.000321	-4.409338	-2.2285	
20	0.505409					
	black 0.188833	0.213062	0.223870	1.128308	1.1855	
45	1.050728					
	hisp -0.037701	-0.055459	-0.049476	1.471019	1.3123	
32	0.892125					
	married -0.025437	-0.026148	-0.027580	1.027952	1.0842	
38	1.054755					

=====

====

SCALING FACTORS

=====

====

Linear Model: APE = β_j (constant partial effect)

Probit Model: Mean($\phi(X'\beta)$) = 0.336235

Logit Model: Mean($\lambda(X'\beta)$) = 0.203070

Ratio: Logit/Probit scaling = 0.6040

```
In [31]: # Create a cleaner comparison (excluding the intercept) and visualize

comparison_clean_all = pd.DataFrame({
    'Variable': ape_ols.index[1:], # Exclude intercept
    'Linear': ape_ols.values[1:],
    'Probit': ape_probit.values[1:],
    'Logit': ape_logit.values[1:]
})

print("*80")
print("AVERAGE PARTIAL EFFECTS: Linear vs Probit vs Logit (Excluding Intercept")
print("*80")
print("\n")
print(comparison_clean_all.to_string(index=False))
print("\n")

# Calculate differences
print("*80")
print("DIFFERENCES IN APE")
print("*80")
print("\nProbit vs Linear:")
print(f" Mean absolute difference: {np.abs(ape_probit.values[1:] - ape_ols.values[1:])}")
print(f" Max absolute difference: {np.abs(ape_probit.values[1:] - ape_ols.values[1:]).max()}\n")

print("\nLogit vs Linear:")
print(f" Mean absolute difference: {np.abs(ape_logit.values[1:] - ape_ols.values[1:])}")
print(f" Max absolute difference: {np.abs(ape_logit.values[1:] - ape_ols.values[1:]).max()}\n")

print("\nLogit vs Probit:")
print(f" Mean absolute difference: {np.abs(ape_logit.values[1:] - ape_probit.values[1:])}")
print(f" Max absolute difference: {np.abs(ape_logit.values[1:] - ape_probit.values[1:]).max()}\n")
```

```
print(f" Max absolute difference: {np.abs(ape_logit.values[1:] - ape_probit
print("\n")
```

```
=====
=====
AVERAGE PARTIAL EFFECTS: Linear vs Probit vs Logit (Excluding Intercept)
=====
```

Variable	Linear	Probit	Logit
train	-0.111703	-0.113173	-0.112330
unem74	0.038693	0.035673	0.039770
unem75	0.015961	0.021389	0.016791
age	0.000043	0.000227	0.000073
educ	0.000144	-0.000636	-0.000321
black	0.188833	0.213062	0.223870
hisp	-0.037701	-0.055459	-0.049476
married	-0.025437	-0.026148	-0.027580

```
=====
=====
DIFFERENCES IN APE
=====
```

Probit vs Linear:

Mean absolute difference: 0.006697
 Max absolute difference: 0.024229

Logit vs Linear:

Mean absolute difference: 0.006498
 Max absolute difference: 0.035037

Logit vs Probit:

Mean absolute difference: 0.003529
 Max absolute difference: 0.010808

Conclusion: Comparing Logit APE with Probit and Linear APE from part (h)

Key Findings:

1. Sign Consistency: All three models (Linear, Probit, Logit) agree on the direction (sign) of the effect for each variable.

2. Magnitude Comparison:

- Linear model APE values are typically the largest in absolute value
- Probit and Logit APE values are smaller because they're scaled by their respective PDFs

- Probit uses $\phi(X'\beta)$, the standard normal PDF
- Logit uses $\Lambda(X'\beta) = \Lambda(X'\beta) \times (1 - \Lambda(X'\beta))$, the logistic PDF

3. Probit vs Logit:

- The APE values from Probit and Logit are very similar
- The logit scaling factor is typically slightly larger than the probit scaling factor
- This is because the logistic distribution has heavier tails than the normal distribution
- In practice, the choice between probit and logit often makes little difference for APE

4. Formula Recap:

- **Linear:** $APE_j = \beta_j$ (constant across all observations)
- **Probit:** $APE_j = \text{mean}(\phi(X'\beta)) \times \beta_j$, where ϕ is the standard normal PDF
- **Logit:** $APE_j = \text{mean}(\Lambda(X'\beta) \times (1 - \Lambda(X'\beta))) \times \beta_j$, where Λ is the logistic CDF

5. Model Selection:

- All three models give similar qualitative conclusions
- Probit and Logit are theoretically superior as they constrain predictions to [0,1]
- Linear is simpler to interpret but can give predictions outside [0,1]
- Choose between Probit/Logit based on assumptions about tail behavior

6. Variable of Interest (train):

- All three models show that training has a negative and statistically significant effect on 1978 unemployment
- The APE is approximately -0.11 to -0.11 across all three models, indicating very consistent results